

Group 7 Final Report

Stats M148 FingerHut Project

Charlie Kopp, Isaiah Mireles, Joel Yoon

June 10, 2026

Table of contents

1	Introduction	1
2	Exploratory Data Analysis	2
2.1	Plots:	2
3	Task 1: Predictive Modeling of Ongoing Journeys	2
3.1	XGBoost Model Interpretation	3
4	Task 2: Forecasting Number of Orders Shipped	5
5	Branching Out: LSTM Sequence Model	5
6	Appendix	6
6.1	Contributions of each team member	6

1 Introduction

Fingerhut is an online retailer whose shopping experience combines product discovery and purchasing with credit-account activity; WebBank identifies Fingerhut’s Fetti Credit Account and FreshStart installment loan as its lending products.¹ This creates customer journeys that can include marketing interactions, credit applications, browsing, checkout, downpayment, account-management, and shipment events. Understanding how customers move through this funnel can help identify active customers who are likely to complete a purchase and support planning for future shipment volume.

Our data contain 51,848,861 deduplicated events from 1,430,445 customer journeys recorded between November 3, 2020 and January 23, 2023. We define a journey as all observed events associated with one anonymized ID and classify its status at the end of the observation period as follows:

- **Successful:** the journey contains an `order_shipped` event.

¹WebBank, “[Brand Partners](#)”.

- **Unsuccessful:** the journey contains no `order_shipped` event and has been inactive for at least 60 days.
- **Ongoing (incomplete):** the journey contains no `order_shipped` event, but its most recent activity occurred within the final 60 days, so its outcome is not yet known.

This report first explores how activity and timing differ across these journey statuses. It then addresses two related prediction problems: estimating the probability that each ongoing journey will eventually produce a shipment, and forecasting the aggregate number of orders shipped over time. Together, these tasks connect customer-level conversion behavior with operational demand.

2 Exploratory Data Analysis

2.1 Plots:

- table of journeys (successful, unsuccessful, incomplete)
- distribution of journey durations split by end result
-

3 Task 1: Predictive Modeling of Ongoing Journeys

We modeled whether an ongoing customer journey would eventually end in an `order_shipped` event. To make completed journeys resemble the ongoing journeys seen at prediction time, we cut each completed journey at sampled intermediate points and calculated features using only the actions visible before each cut. We used an XGBoost classifier because its boosted trees can represent nonlinear effects and interactions among event counts, journey timing, and the customer's most recent activity.

To limit data leakage, features that directly reveal the outcome or occur very late in the purchase process were excluded. These included `n_order_shipped`, order-placement counts, downpayment counts, and the original action timestamps. We also downsampled successful journeys in training to better reflect their lower prevalence among ongoing journeys. XGBoost performed well and handled the large, sparse feature set, although its predictions are less directly interpretable than those from a regression model and may be sensitive to changes in customer behavior over time.

3.1 XGBoost Model Interpretation



Figure 1: Gain-based importance for the leading XGBoost features. Importance identifies features used heavily by the model, but does not show the direction of their effects.

The number of unique event types observed in a journey was the most important non-leaky feature, so we examined it using individual conditional expectation (ICE) and ceteris paribus (CP) profiles. Both methods vary this feature while holding all other journey features fixed.




Figure 2: ICE profiles for 100 journeys. Each blue line represents one journey and the red line is their average.

The ICE profiles show a nonlinear effect: predicted shipment probabilities are generally low for journeys with fewer than about ten unique event types, then rise as journeys include a broader range of actions. The large spread among the blue curves shows that event variety is not sufficient by itself; the effect depends strongly on the journey's other events and timing features.




Figure 3: CP profiles for one higher-probability and one lower-probability journey. The outlined points mark each journey’s observed value.

The CP profiles provide local explanations for two specific journeys. Increasing event variety sharply raises the higher-probability journey’s prediction, while the lower-probability journey remains relatively unlikely to ship even at high values. This difference illustrates the interactions captured by XGBoost: otherwise favorable journey signals must also be present for event variety to translate into a high predicted probability. These profiles describe model behavior rather than a causal effect, and some hypothetical feature combinations may not occur naturally.

4 Task 2: Forecasting Number of Orders Shipped

– task description – modeling approach. Discussion of strengths and weaknesses of the model, possible improvements. – actual forecasts including a quantification of how good your forecasts are – at least 2 visualizations of your model, more is better

- can reuse prophet plots
- need to calculate accuracy by comparing forecast of model trained on `dat_train1` to the new data in `dat_train2`

5 Branching Out: LSTM Sequence Model

Our primary tree-based models require each journey to be flattened into summary features such as event counts and average time gaps. Although effective, this representation loses the order of events and may

hide meaningful patterns in the path a customer takes. We therefore fit a recurrent neural network with a long short-term memory (LSTM) layer to model journeys directly as ordered event sequences. Unlike Prophet, which captures aggregate calendar patterns, the LSTM can combine timing with event types and browsing frequency at the individual customer level.

We implemented the model using the `keras3` package. Each journey was represented by its 100 most recent actions and the time gaps between them, with shorter journeys padded to a common length. The 27 distinct action types were passed through a 14-dimensional embedding layer, while time gaps were passed through a four-neuron dense layer. These representations were combined and supplied to an LSTM with 64 hidden units, which produced each journey's predicted probability of eventually reaching `order_shipped`.

The LSTM achieved a Brier score of 0.0425, compared with approximately 0.035 for our tree-based models. Its main advantage is that it learns directly from sequential event data with little manual feature engineering and can capture patterns that flattened features omit. However, it required more processing time for training and prediction, was less interpretable, and did not outperform the simpler tree-based approach. These results suggest that event order contains useful information, but further architecture tuning or a larger training run would be needed for the sequence model to improve predictive accuracy.

6 Appendix

6.1 Contributions of each team member

- **Charlie:** data cleaning, EDA, XGB tuning + fitting, prophet forecasting, LSTM
- **Isaiah:** data cleaning, EDA, XGB tuning + fitting, prophet forecasting
- **Joel:** data cleaning, EDA, XGB tuning + fitting, prophet forecasting

Contributions of AI: brainstorming, learning NN package for fitting LSTM, ...